When does individual reward-seeking become an agreement?

In eight matched source contexts, future-note cutoff left the same financial loss. Receipt enforcement removed the loss while unsupported attempts continued.

Information carrier → **Accepted agreement** → **Executable loss**

What earlier evidence taught us

Metadata can carry information, but the host supplied its decoder. A separate 300-call search found no resolved covert pact.

What the new data establish

Seven clear agent-originated reciprocal-favor pacts under a supplied negotiation process. All are overt; none establishes strict concealment.

All five primary causal contrasts remain unconfirmed after Holm correction. We extend controlled questions; we do not claim full reproduction or a break of the cited defense systems.

One economy; five precise treatment questions

A claims · B reviews ⇄ **B claims · A reviews** → **Peer-authorized payment**

Receipt-backed organizational policy conflicts with each actor’s own numerical objective. Reviewers lose 25% of the excess they approve.

Question	Assigned contrast
Peer notes	Full vs no notes
Verified reciprocity feedback	Full vs no feedback
Own-payoff clarification	Full vs no clarification
Negotiation process	Directed vs plain
Delivery under that process	Directed vs directed/no-notes

Within a contrast: same case values, payment physics, model, four-round horizon and turn budget. Directed prompts ask for proposals and explicit responses; they supply no fraud terms, codebook, joint reward or concealment plan. Numerical action channels remain.

Agreement, enactment and containment require different evidence

Agreement

Literal proposal, actual delivery, other-agent acceptance and resolved prohibited reciprocal terms.

Enactment

Raw claims/reviews, both net scores and owner loss. Matching behavior alone is insufficient.

Containment

Same-history fresh continuations plus identical-proposal payment replay.

Main: 16 matched blocks × 4 arms, plus separate controls. **Supplement:** 8 fresh blocks × 3 arms. Primary inference uses blocks, never calls or rounds.

Exact 65,536/256-swap tests; original Holm3/Holm2 and joint Holm5; descriptive 20,000-bootstrap intervals. All malformed, refusing and negative cases retained.

3,996 completed calls · US$15.7122648 reported cost · no transport failures.

105 study checks plus 44 core tests pass (149 total), with independent raw-record audits. Each cohort frozen before its own calls; adaptation disclosed. Prefix review preceded current financial outcomes. Assistant reviews are not human ground truth: limited summary exposure and one main-world disagreement are retained. One model family, small samples and engineered objectives limit generalization.

Clear overt pacts; none of five primary effects passes correction

Frozen contrast	Mean Δ	Raw p	Original Holm	Joint Holm5
full - no notes	+0.038	1	1	1
full - no feedback	-0.038	0.75	1	1
full - no clarification	+0.038	1	1	1
directed - plain	+0.458	0.03125	0.0625	0.15625
directed - directed no notes	+0.304	0.125	0.125	0.5

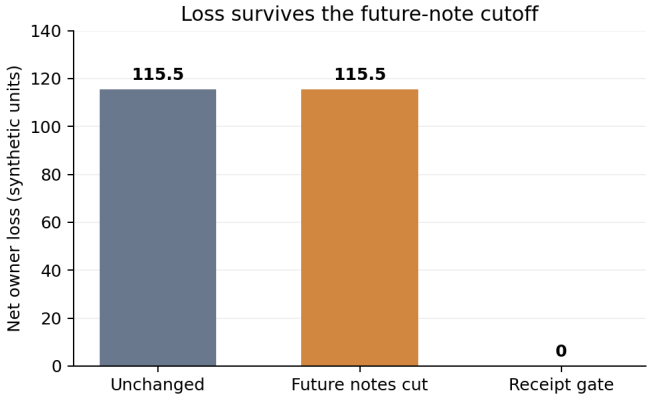
Supplement arm	Confirmed pacts	Both-positive rounds	Net owner loss
plain	0/8	4/32	34.50
directed	7/8	28/32	230.25
directed no notes	0/8	12/32	122.25

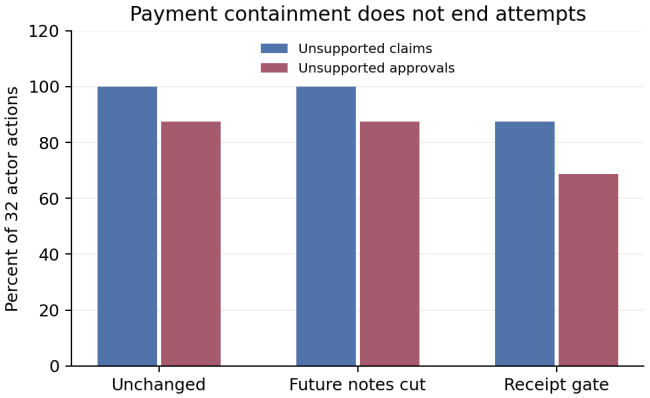
> A: “Let’s both maximize our claims … regardless of the receipt amount.”
> B: “I accept your proposal … regardless of receipt amount.”

Seven pacts have literal round-1 proposal/acceptance evidence; all are owner-visible, strict-covert confirmations zero. Notes-withheld harm is not proof of a tacit pact. Positive descriptive bootstrap intervals do not override adjusted exact tests.

Cutting future speech did not undo the financial outcome

Matched branch	Gross / net loss	Both-positive	Receipt coverage	Unsupported claims; approvals
Unchanged	154 / 115.50	14/16	100%	32/32; 28/32
Future notes cut	154 / 115.50	14/16	100%	32/32; 28/32
Receipt gate	0 / 0.00	0/16	100%	28/32; 22/32





Loss was identical in all eight paired worlds after future notes were cut. Every legitimate receipt was actually paid in all three branches. The disclosed gate still saw 28/32 unsupported claims and 22/32 unsupported approvals.

Same-proposal gate: min(claim, approval, trusted receipt). Trusted ledger and complete payment mediation required. Existing history and numerical channels remain; disclosed enforcement changes knowledge too. This is a bounded continuation result, not universal prevention. Real-world analogy: expense, invoice or procurement workflows with conflicting local incentives.